



# The Edge of Machine Learning

Resource-efficient ML in 2 KB RAM  
for the Internet of Things

Aditya Kusupati & Don Dennis  
Microsoft Research India



D. Dennis, S. Gopinath, P. Jain, A. Kusupati, N. Natarajan, S. Patil, R. Sharma, H. Simhadri & M. Varma

# Resource-constrained IoT Devices

---

## Freescal KL03 microcontroller

ARM® Cortex®-M0+ processor

48 MHz

32 KB Flash, 8KB boot ROM, 2 KB RAM

35uA/MHz low-power active mode

1 uA sleep mode



2mm x 1.6mm

ARM Cortex M0+ at 48 MHz & 35  $\mu$ A/MHz  
with 2 KB RAM & 32 KB read only Flash


# Edge Machine Learning – Objectives

---

- To build a library of machine learning algorithms
  - Which can be trained in the cloud
  - But which will run on tiny IoT devices



ARM Cortex M0+






# Microsoft's EdgeML Library

---


- Compact tree, kNN and RNN algorithms for classification, regression, ranking, time series *etc.*,




Bonsai



ProtoNN



EMI-RNN



FastGRNN

<https://github.com/Microsoft/EdgeML>

# Time Series

- Time series are the most frequently occurring types of signals found in the IoT domain




“Hey, Cortana”



Hey,

Cor


tana



Sprinklers



Soil moisture during a day



Smart farm



# FastGRNN




A Fast, Accurate, Stable & Tiny  
(Kilobyte Sized) Gated RNN



A. Kusupati (MSRI), M. Singh (IITD), K.  
Bhatia (Berkeley), A. Kumar (Berkeley),  
P. Jain (MSRI) & M. Varma (MSRI)

# Recurrent Neural Networks (RNNs)

- State-of-the-art for analyzing sequences & time series
- Training is unstable due to exploding & vanishing gradients

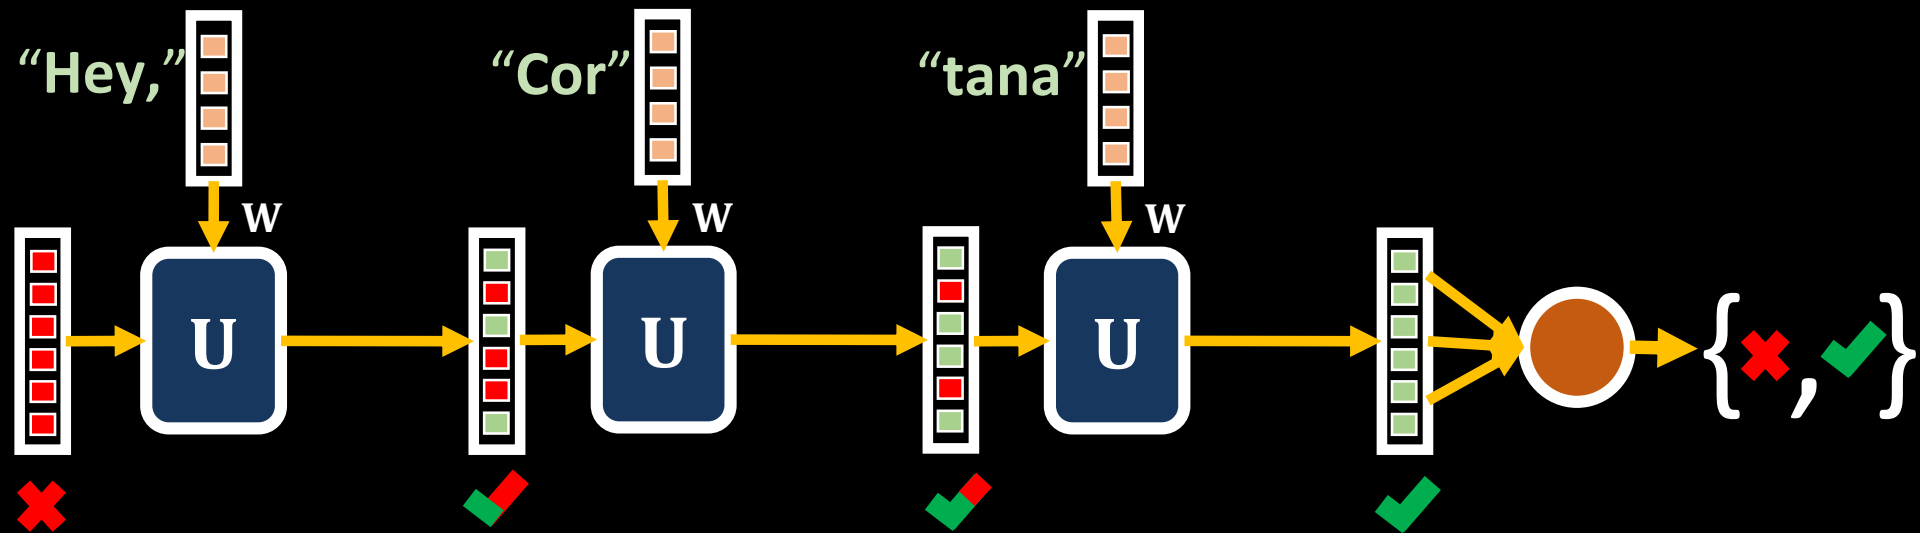


$$\mathbf{h}_t = \sigma(\mathbf{W}\mathbf{x}_t + \mathbf{U}\mathbf{h}_{t-1} + \mathbf{b})$$



# Recurrent Neural Networks (RNNs)


- State-of-the-art for analyzing sequences & time series
- Training is unstable due to exploding & vanishing gradients



$$\nabla = f(\dots, \mathbf{U}^T = \mathbf{Q} \begin{bmatrix} 1.2^{100} & & \\ & \ddots & \\ & & 0.8^{100} \end{bmatrix} \mathbf{Q}^T, \dots)$$

# Unitary RNNs – uRNN, SpectralRNN, ...

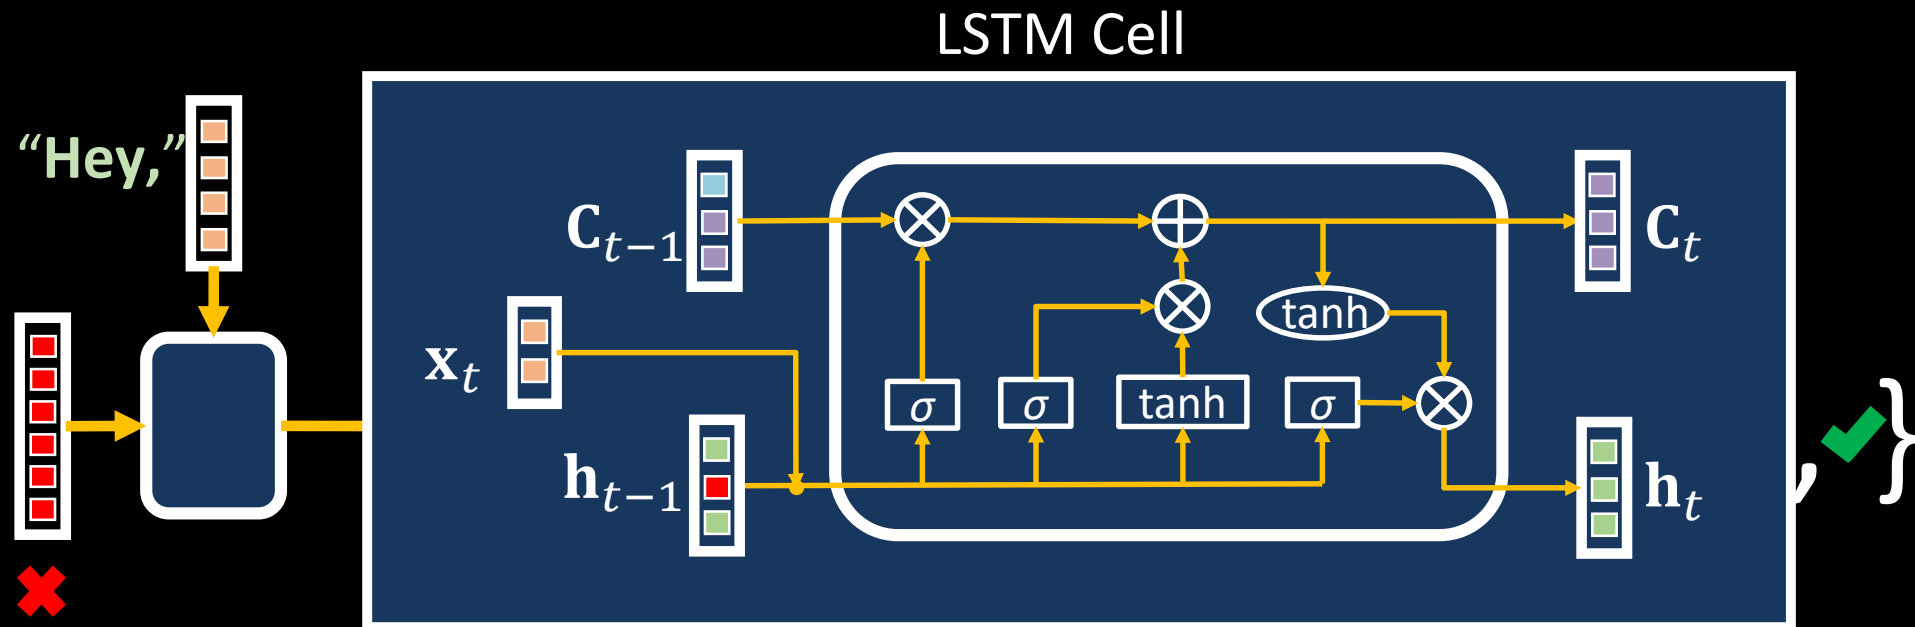
- Unitary RNNs force all the eigenvalues of  $\mathbf{U}$  to be  $\approx 1$
- Unfortunately, they are expensive to train & lack accuracy



$$\nabla = f(\dots, \mathbf{U}^T = \mathbf{Q} \begin{bmatrix} 1^T & & \\ & \ddots & \\ & & 1^T \end{bmatrix} \mathbf{Q}^T, \dots)$$


# Gated RNNs – LSTM, GRU, ...

- Add extra parameters to stabilize training
- Have increased prediction costs on IoT microcontrollers
- Have intuitive explanations but lack formal guarantees



# FastRNN

- Provably stable training with 2 additional scalars
- Accuracy: RNN  $\ll$  Unitary RNNs  $<$  FastRNN  $<$  Gated RNNs



$$\begin{aligned}\widetilde{\mathbf{h}}_t &= \sigma(\mathbf{W}\mathbf{x}_t + \mathbf{U}\mathbf{h}_{t-1} + \mathbf{b}) \\ \mathbf{h}_t &= \alpha\widetilde{\mathbf{h}}_t + \beta\mathbf{h}_{t-1}\end{aligned}$$

# RNN gradients

---

$$\mathbf{h}_t = \sigma(\mathbf{W}\mathbf{x}_t + \mathbf{U}\mathbf{h}_{t-1} + \mathbf{b})$$

$$\frac{\partial L}{\partial \mathbf{U}} = \sum_{t=0}^T \mathbf{D}_t \left( \prod_{k=t}^{T-1} \mathbf{U}^\top \mathbf{D}_{k+1} \right) (\nabla_{\mathbf{h}_T} L) \mathbf{h}_{t-1}^\top$$

$$\frac{\partial L}{\partial \mathbf{W}} = \sum_{t=0}^T \mathbf{D}_t \left( \prod_{k=t}^{T-1} \mathbf{U}^\top \mathbf{D}_{k+1} \right) (\nabla_{\mathbf{h}_T} L) \mathbf{x}_t^\top$$

$$\mathbf{D}_k = \text{grad}(\mathbf{h}_k) \quad \mathbf{D}_k = \text{diag}(\sigma'(\mathbf{W}\mathbf{x}_k + \mathbf{U}\mathbf{h}_{k-1} + \mathbf{b}))$$



# FastRNN gradients

$$\begin{aligned}\tilde{\mathbf{h}}_t &= \sigma(\mathbf{W}\mathbf{x}_t + \mathbf{U}\mathbf{h}_{t-1} + \mathbf{b}) \\ \mathbf{h}_t &= \alpha\tilde{\mathbf{h}}_t + \beta\mathbf{h}_{t-1}\end{aligned}$$

$$\frac{\partial L}{\partial \mathbf{U}} = \alpha \sum_{t=0}^T \mathbf{D}_t \left( \prod_{k=t}^{T-1} (\alpha \mathbf{U}^\top \mathbf{D}_{k+1} + \beta \mathbf{I}) \right) (\nabla_{\mathbf{h}_T} L) \mathbf{h}_{t-1}^\top$$

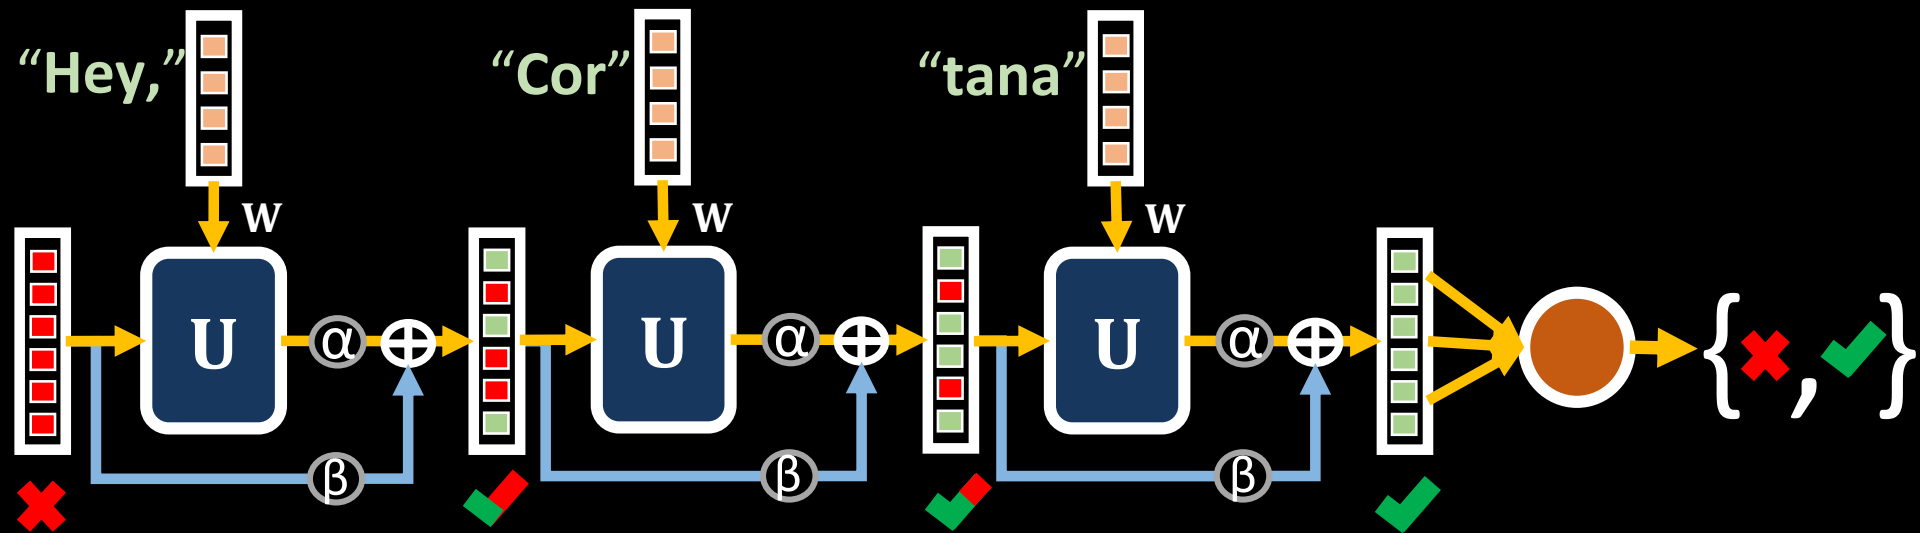
$$\frac{\partial L}{\partial \mathbf{W}} = \alpha \sum_{t=0}^T \mathbf{D}_t \left( \prod_{k=t}^{T-1} (\alpha \mathbf{U}^\top \mathbf{D}_{k+1} + \beta \mathbf{I}) \right) (\nabla_{\mathbf{h}_T} L) \mathbf{x}_t^\top$$

$$\mathbf{D}_k = \text{grad}(\tilde{\mathbf{h}}_k)$$

$$\mathbf{D}_k = \text{diag}(\sigma'(\mathbf{W}\mathbf{x}_k + \mathbf{U}\mathbf{h}_{k-1} + \mathbf{b}))$$

# FastRNN

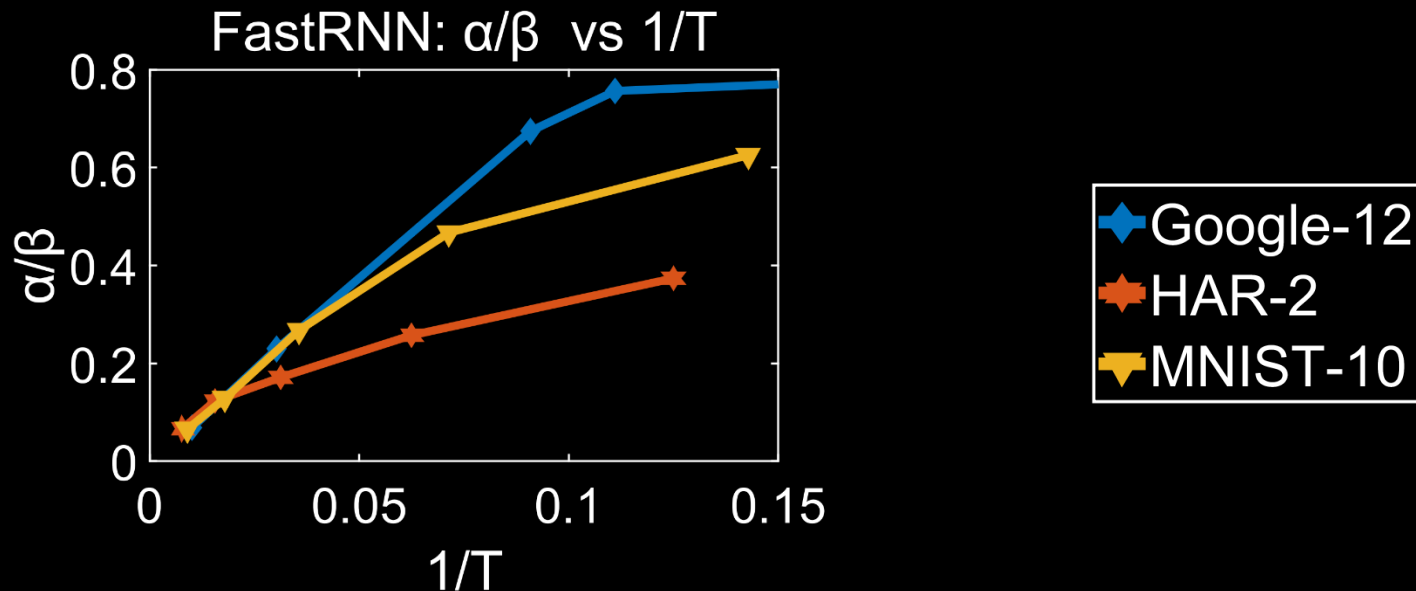
- Provably stable training with 2 additional scalars
- Accuracy: RNN  $\ll$  Unitary RNNs  $<$  FastRNN  $<$  Gated RNNs



$$\nabla = f(\dots, (\alpha \mathbf{U} \mathbf{D} + \beta \mathbf{I})^T = \mathbf{Q} \begin{bmatrix} (\beta + \alpha \|\mathbf{U} \mathbf{D}\|)^T \\ \vdots \\ (\beta - \alpha \|\mathbf{U} \mathbf{D}\|)^T \end{bmatrix} \mathbf{Q}^T, \dots)$$

# Theoretical Analysis of FastRNN

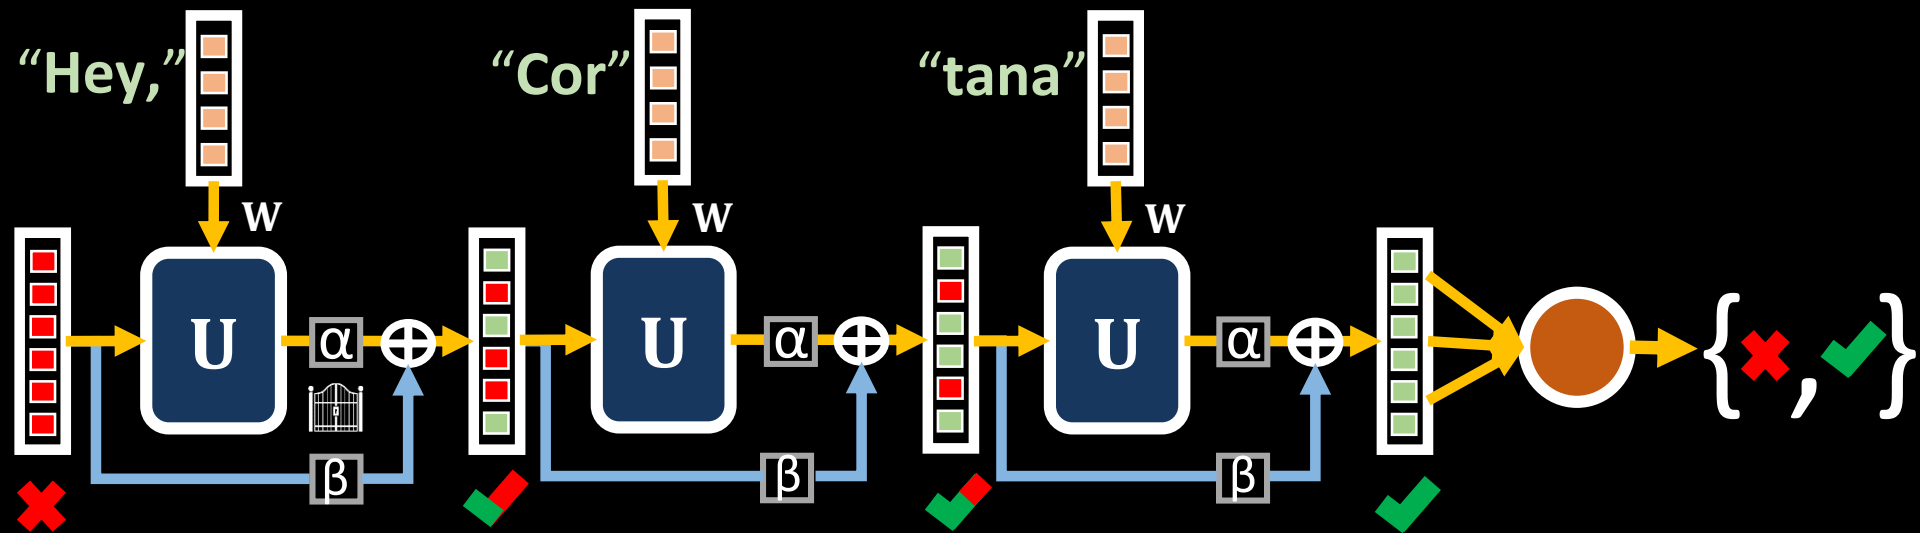
- Convergence & generalization error bounds
  - FastRNN is independent of  $T$  when  $\alpha/\beta \approx O(1/T)$
  - A similar analysis reveals that an RNN is exponential in  $T$



- FastRNN converges to within  $\epsilon$  of a stationary point in  $O(1/\epsilon^2)$  SGD iterations and its generalization error is bounded by  $O(\frac{1}{\sqrt{n}})$

# FastGRNN

- Extend  $\alpha$  &  $\beta$  from scalars to vector gates
- Accuracy: RNN  $\ll$  Unitary RNNs  $<$  Gated RNNs  $\approx$  FastGRNN



$$\beta_t = \sigma_\beta(\mathbf{W}\mathbf{x}_t + \mathbf{U}\mathbf{h}_{t-1} + \mathbf{b}_\beta); \tilde{\mathbf{h}}_t = \sigma_h(\mathbf{W}\mathbf{x}_t + \mathbf{U}\mathbf{h}_{t-1} + \mathbf{b}_h)$$

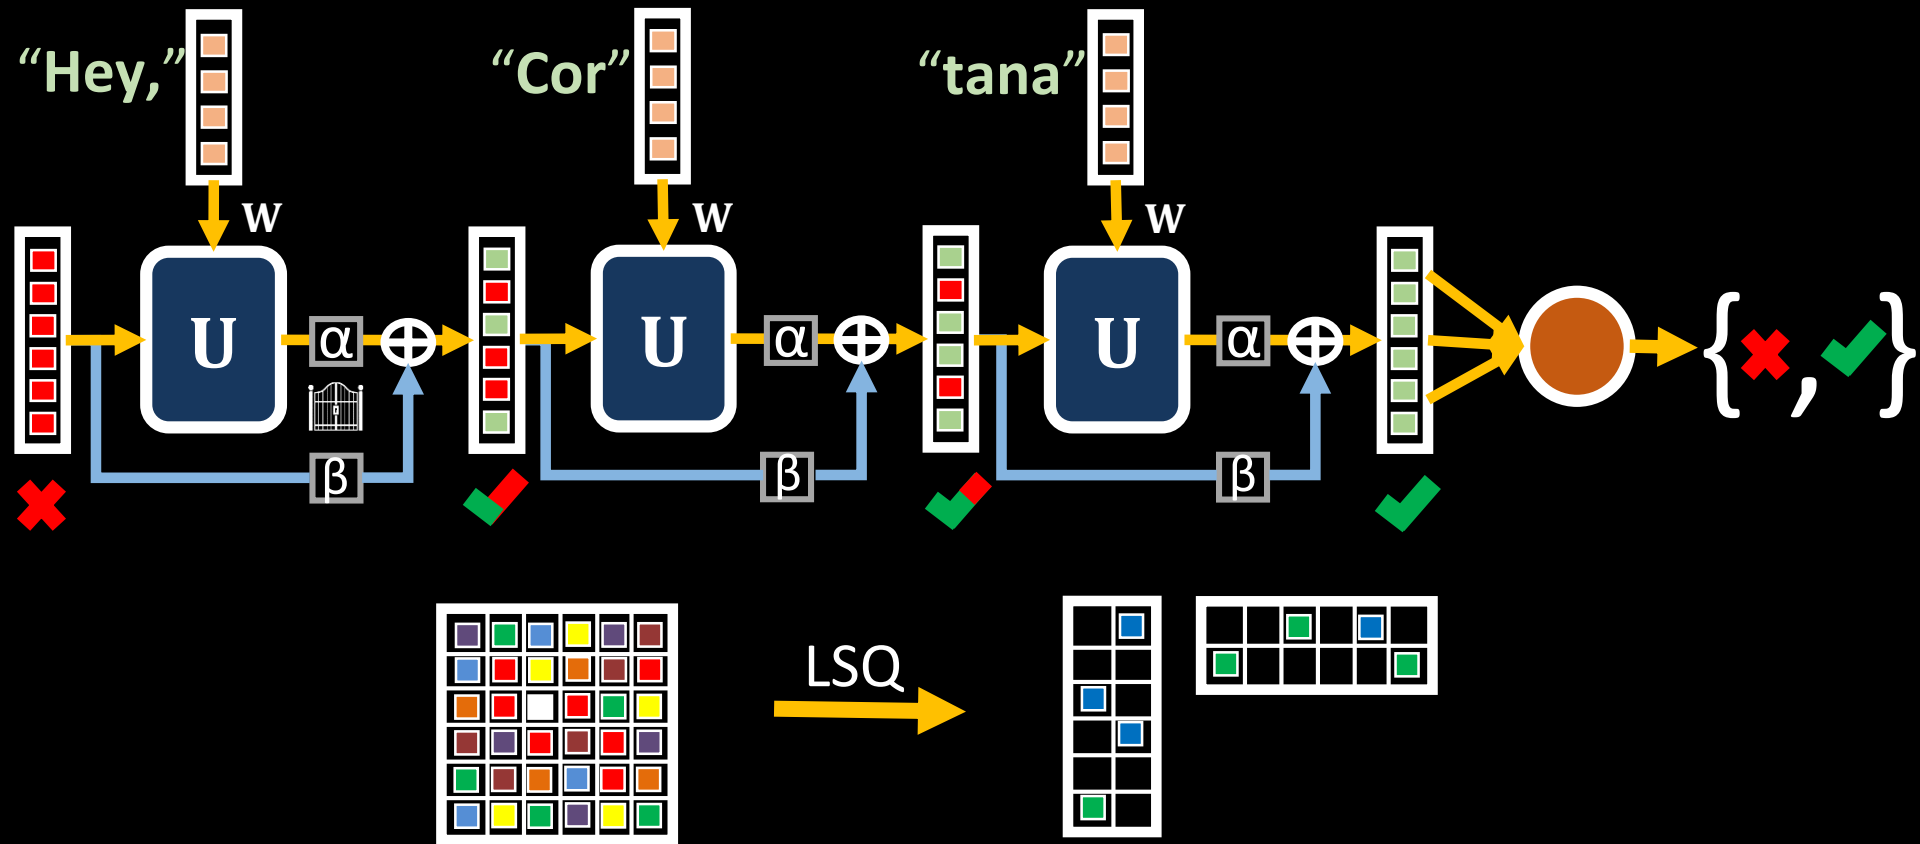
$$\alpha_t \approx \mathbf{1} - \beta_t; \quad \mathbf{h}_t = \alpha_t \odot \tilde{\mathbf{h}}_t + \beta_t \odot \mathbf{h}_{t-1}$$





# FastGRNN

- Make  $U$  and  $W$  low-rank (L), sparse (S) and quantized (Q)
- Model Size: FastGRNN  $\ll$  RNN  $\approx$  Unitary RNNs  $<$  Gated RNNs




# Dataset Statistics

	Dataset	# Train	# Features	# Time Steps	# Test
Speech	Google-12	22,246	3,168	99	3,081
	Google-30	51,088	3,168	99	6,835
	Wakeword-2	195,800	5,184	162	83,915
NLP	Yelp-5	500,000	38,400	300	500,000
	PTB-10000	929,589	---	300	82,430
Activity	HAR-2	7,352	1,152	128	2,947
	DSA-19	4,560	5,625	125	4,560
Image	Pixel-MNIST-10	60,000	784	784	10,000


# Comparison to Gated Architectures

- Uncompressed FastGRNN is as accurate as a GRU/LSTM
- FastGRNN is almost as accurate as a GRU/LSTM (within 1%)
- FastGRNN is 20-80x smaller than a GRU/LSTM



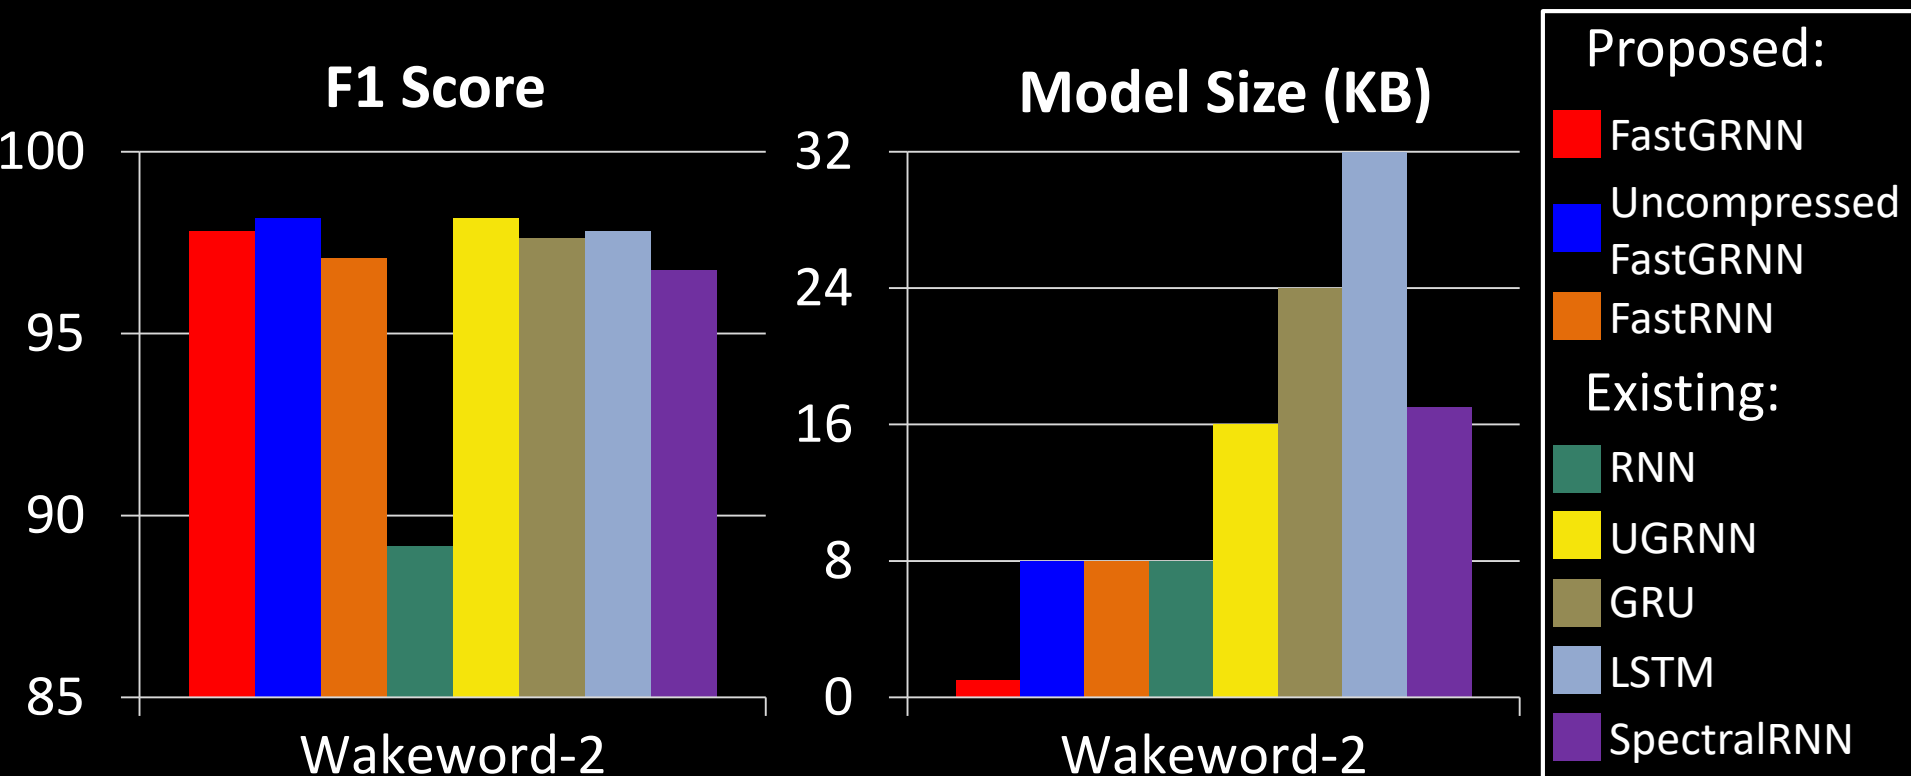
# Comparison to Unitary Architectures

- FastRNN outperforms all unitary RNNs on most datasets
- FastGRNN can be 3-5% more accurate
- FastGRNN can be 45x-200x smaller



# Recognizing “Hey, Cortana” in 1 KB


- Uncompressed FastGRNN outperforms state-of-the-art RNNs
- FastGRNN matches state-of-the-art RNN accuracies





# Prediction on Edge Devices

- None of the other RNNs fit on an Arduino Uno
- FastGRNN can be 25-132x faster at prediction on the MKR1K





# EMIRNN

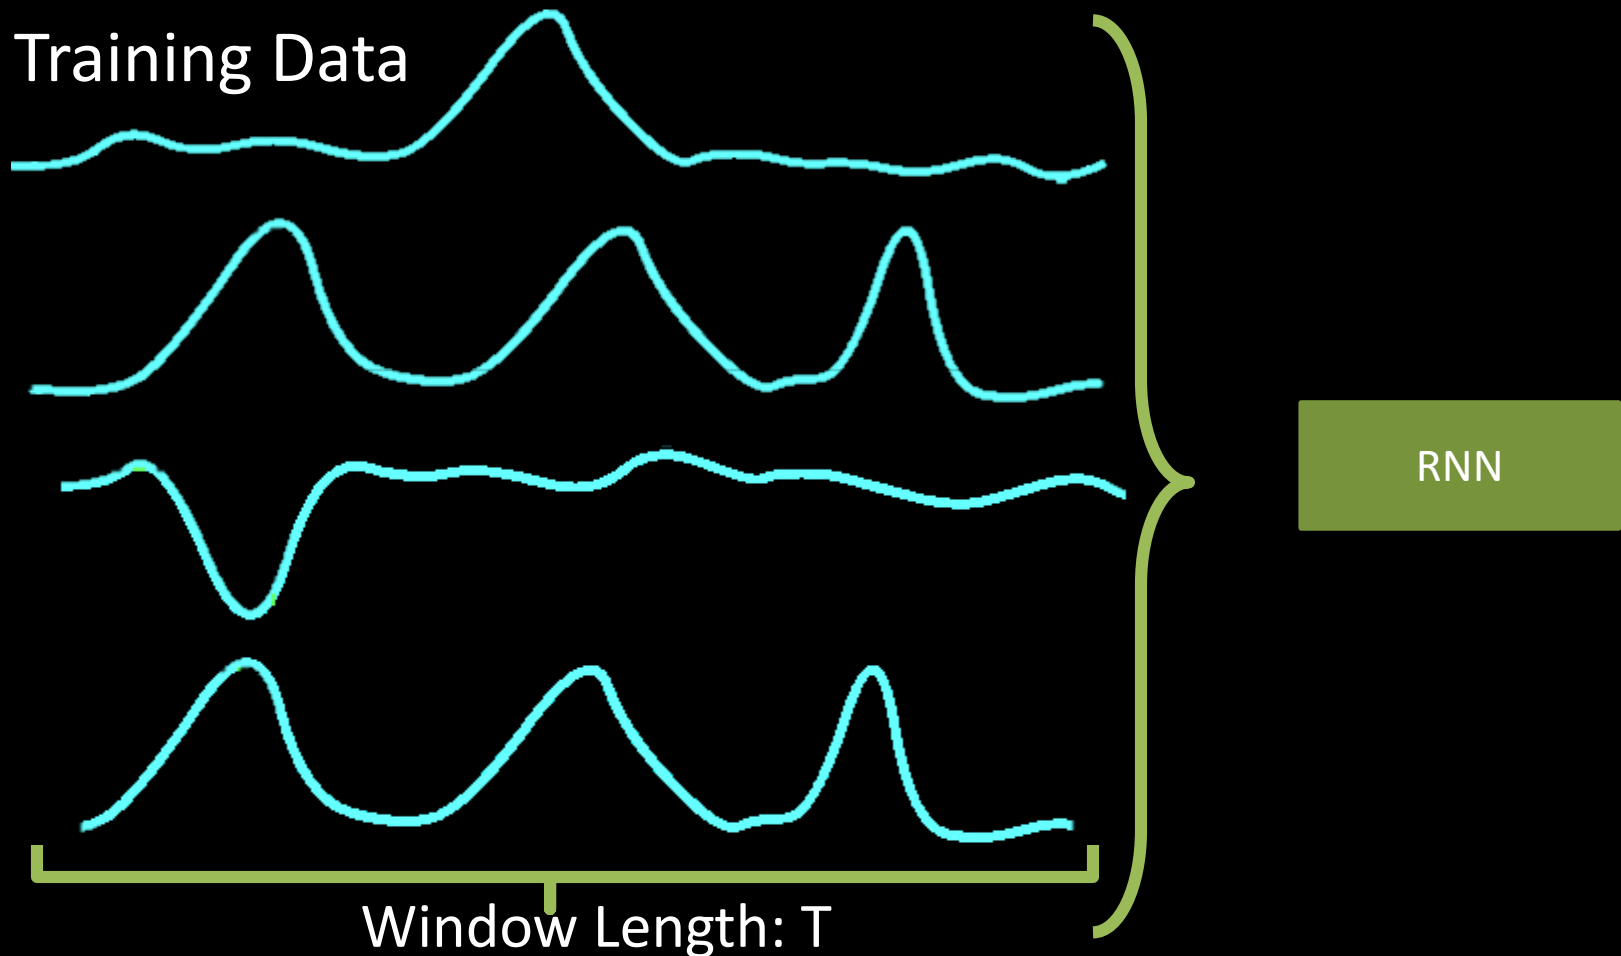


Simultaneous segmentation and  
early prediction in RNNs




Don Dennis (MSRI), Chirag P (MSRI),  
Harsha Simhadri (MSRI), P. Jain (MSRI)

# Time-series Classification: Training



# Time-series Analysis: Sliding Windows



Prediction Cost per Window:  $O(\text{Window Size})$


# Time-series Analysis: Sliding Windows

---

- Complicated RNN cell updates
- Running time  $O(T)$
- Information reuse across windows




# Coarse Training Data



Typically  $k \ll T$ , i.e., actual signature of event is tiny

- Audio clips: 2-5secs but “Hey Cortana” typically spoken in <1sec
- *Unnecessarily large  $T$*  --- longer prediction time, lag
- Predictors must recognize signatures with different offsets
  - *requires larger* predictors.


# Smaller Windows?



- Only a few true positives: **several false positives!**
- Issue: apriori location of true positives unknown

# EMI-RNN: Approach

---

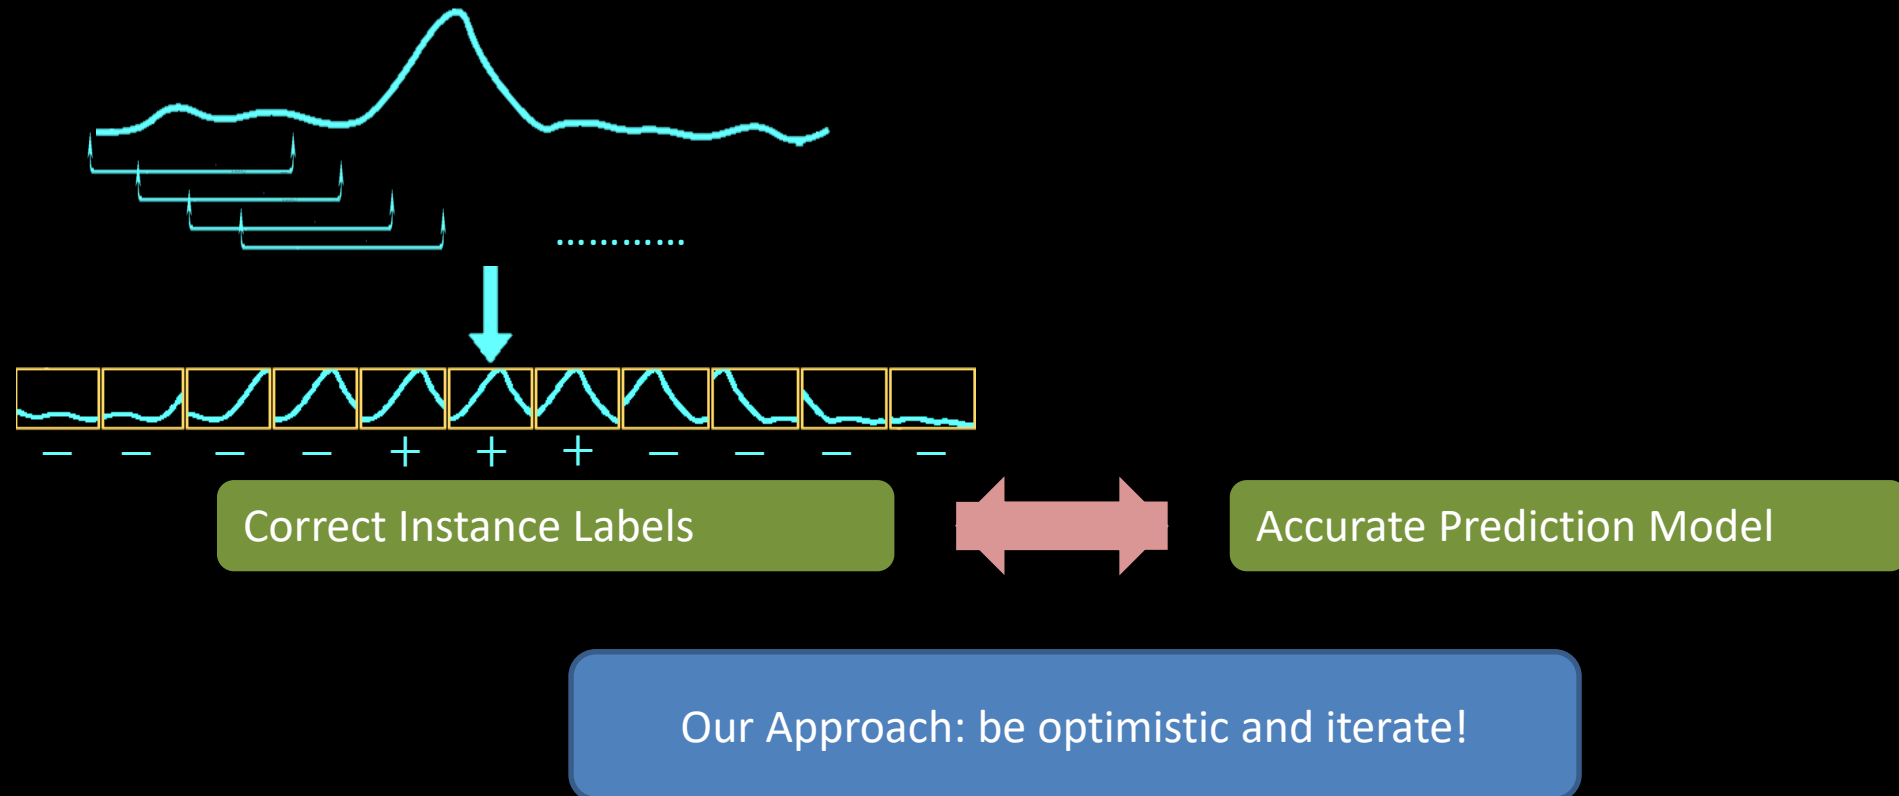


Exploit temporal locality with MIL/Robust learning techniques

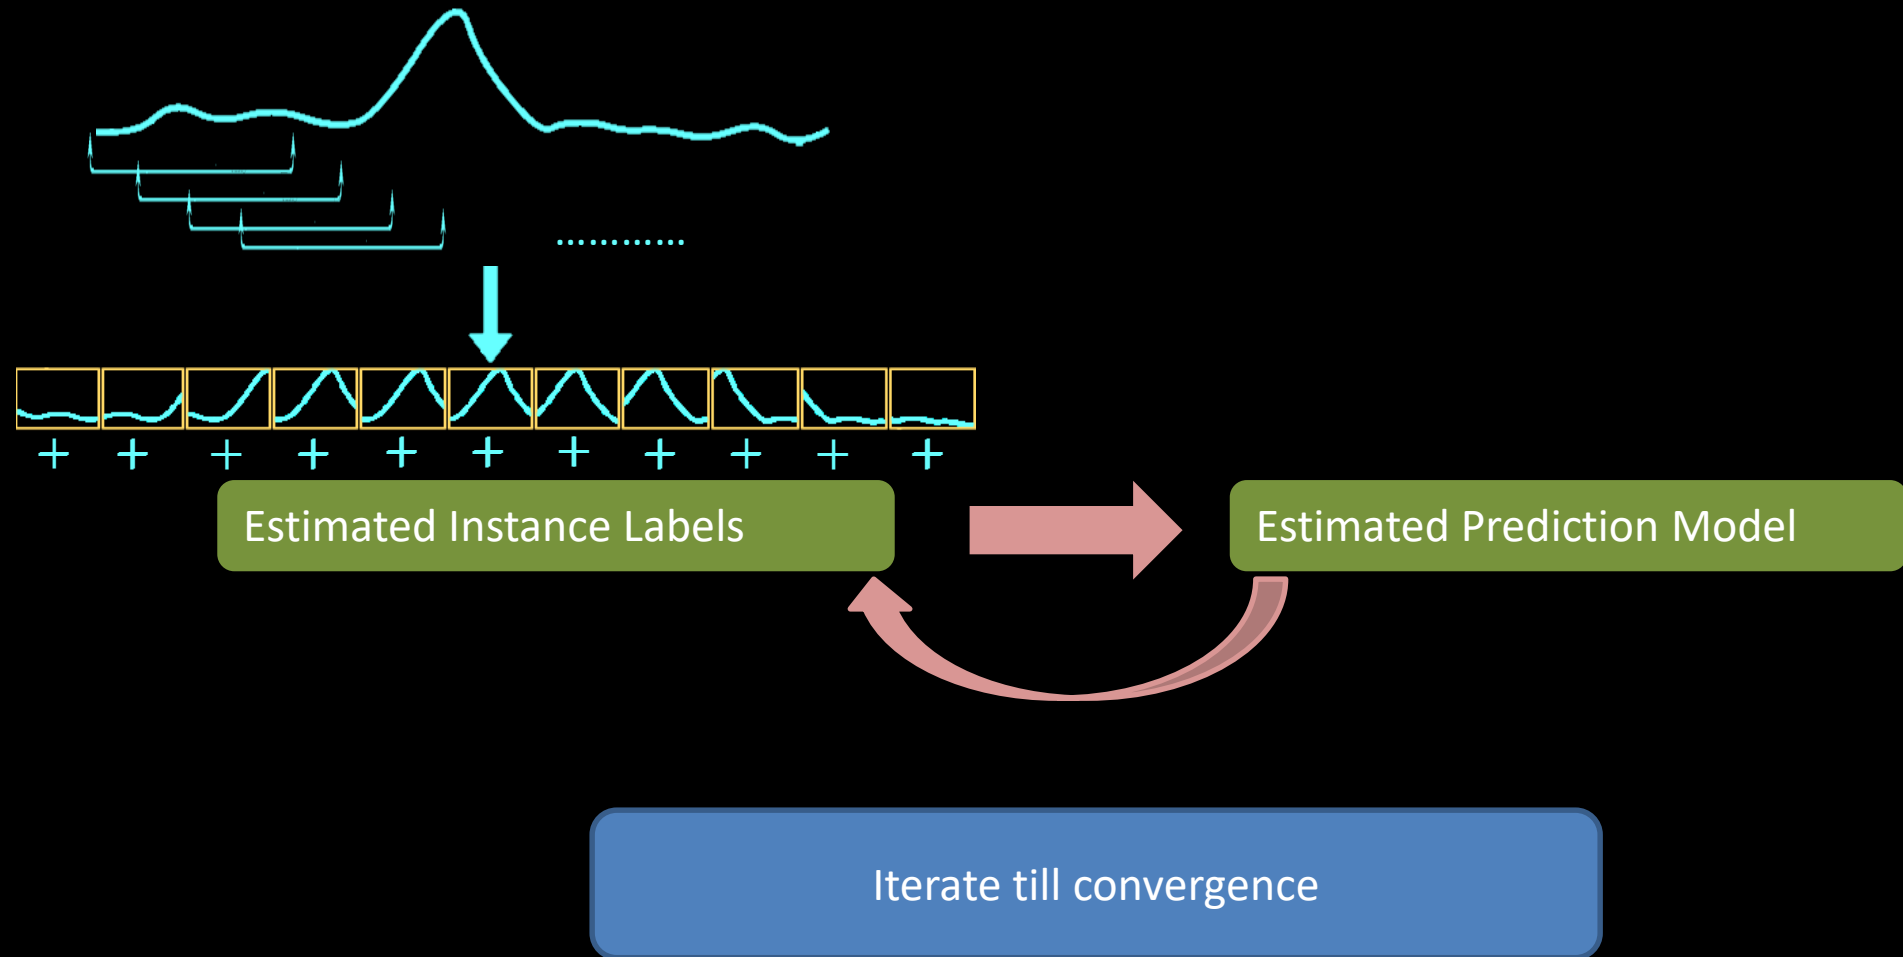
Property 1: Positive instances are clustered together.

Property 2: Number of positive instances can be estimated.

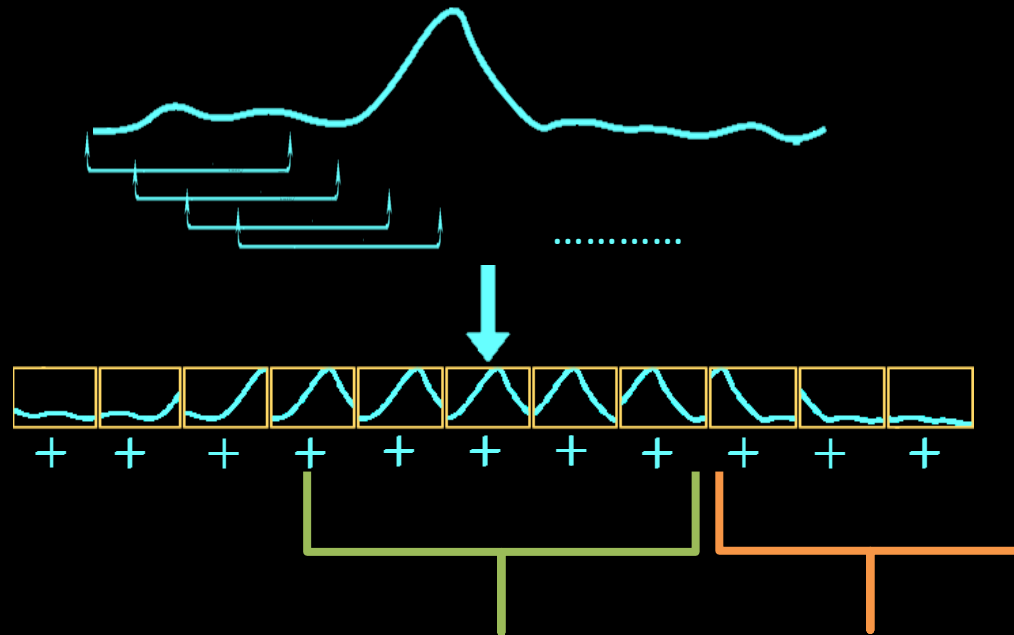
# EMI-RNN: Chicken & Egg Problem



# EMI-RNN: Algorithm



# EMI-RNN: Algorithm



Correct Labels:  
Unique to this class

Incorrect Labels:  
Found in multiple classes

# EMI-RNN: Analysis?

- Optimizing:

$$\min_{W, \hat{y}_{ij}} \sum_{i=1}^n \sum_{j=1}^m \text{loss}(\hat{y}_{ij}, f(z_{ij}, W))$$


$$\text{s.t. } \sum_j \hat{y}_{ij} = k, \quad i: \text{positive}$$

$\hat{y}_{ij}$  satisfies temporal locality

$$\hat{y}_{ij} \in \{0,1\}, i: \text{positive}, \quad \hat{y}_{ij} = 0, i: \text{negative}$$

- $f$ : prediction function (LSTM/FastGRNN)
- Algorithm: alternating between  $\hat{y}_{ij}$  and  $W$

# EMI-RNN: Analysis?




Alternating minimization for Non-convex optimization problem

- Need not converge in general!

*Theorem:* In  $\log n$  iterations, the true positive set

$$S_* = \{(i, j), \hat{y}_{ij} = +1\}$$

will be recovered exactly, with high probability.




Positives need not be i.i.d.

Non-homogenous setting: first such result in literature




# Early Prediction?




- Existing work:
  - Assumes pretrained classifier and uses secondary classifiers
  - Template matching approaches
  - Separate policy for early classification

# Early Prediction?

Positive Datapoint



Negative Datapoint



## Our Approach


**Inference:** Predict at each step – stop as soon as prediction confidence is high.

**Training:** Incentivize early prediction by rewarding correct and early detections.

# Advantages

- Obvious: Speed
- Prediction accuracy
- Early prediction

# EMI-RNN: Empirical Results



# EMI-RNN: Empirical Results

---

Dataset	Accuracy Gain (over Baseline LSTMs)	Prediction Time Reduction
HAR	0.8%	8x
Sports	2.0%	8x
Google	1.5%	8x
Interactive Cane	1.0%	45x

# EMI-RNN: Empirical Results

---

Dataset	Accuracy Gain (over Baseline LSTMs)	Prediction Time Reduction	Memory Savings
HAR	0.7%	11x	4x
Sports	2.0%	12x	4x
Google	1.5%	8x	4x
Interactive Cane	0.9%	72x	36x

# Time-series Analysis: Conclusions

---

- Complicated RNN cell updates: **FastGRNN**
- Running time  $O(T)$ : **EMI-FastGRNN**
- Information reuse across windows:  
**Shallow Recurrent Networks**

# EdgeML repository

---

- Code release
  - EdgeML: 70K page views, 1017 clones, 569 stars
  - Bonsai and ProtoNN released as TLC Beta
- Used extensively in MSR India IoT Summer School
  - Automated voice feedback system
  - Radar-based poacher detection and SONYC
  - Predictive maintenance for solar panels
- <https://github.com/Microsoft/EdgeML>



Thank You